

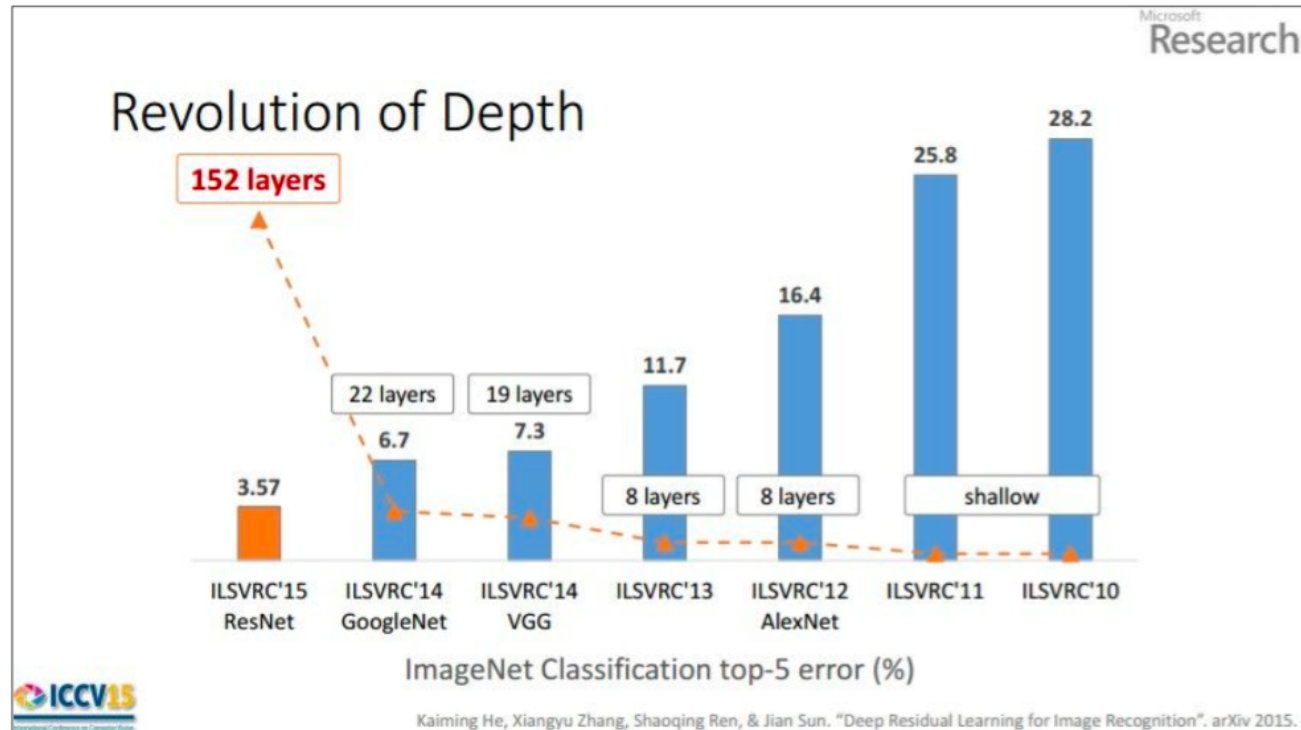
TTIC 31230, Fundamentals of Deep Learning

David McAllester, Winter 2020

Convolutional Neural Networks (CNNs)

Imagenet Classification

1000 kinds of objects.



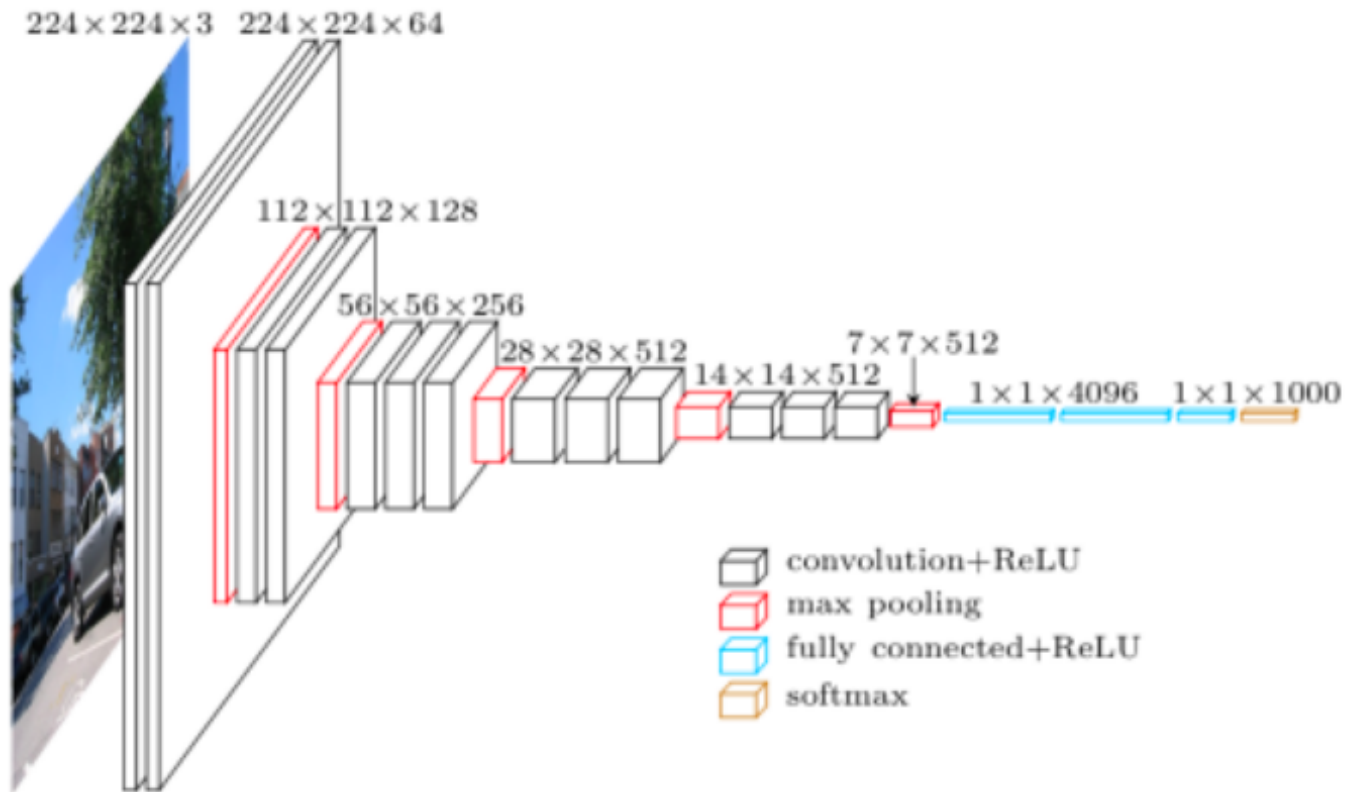
(slide from Kaiming He's recent presentation)

2016 is 3.0%, is 2017 2.25%


SOTA as of January 2020 is 1.3%

What is a CNN?

VGG, Zisserman, 2014




A Convolution Layer



Each box is a tensor $L_\ell[b, x, y, i]$

For a convolution layer, each $L_{\ell+1}[b, x, y, j]$ is the output of a single linear threshold unit computed from $L_\ell[b, x, y, i]$.

A Convolution Layer



$$W[\Delta x, \Delta y, i, j]$$

$$L_{\ell}[b, x, y, i]$$

$$L_{\ell+1}[b, x, y, j]$$

River Trail Documentation

$$L_{\ell+1}[b, x, y, j]$$

$$= \sigma \left(\left(\sum_{\Delta x, \Delta y, i} W[\Delta x, \Delta y, i, j] L_{\ell}[b, x + \Delta x, y + \Delta y, i] \right) - B[j] \right)$$

Many “Neurons” (Linear Threshold Units)

Each $L_{\ell+1}[b, x, y, j]$ is the output of a single linear threshold unit.

$$L_{\ell+1}[b, x, y, j] \\ = \sigma \left(\left(\sum_{\Delta x, \Delta y, i} W[\Delta x, \Delta y, i, j] L_{\ell}[b, x + \Delta x, y + \Delta y, i] \right) - B[j] \right)$$

2D CNN in PyTorch

`conv2d(input, weight, bias, stride, padding, dilation, groups)`

input tensor (minibatch,in-channels,iH,iW)

weight filters (out-channels, in-channels/groups, in-channels,kH,kW)

bias tensor (out-channels) . Default: None


stride Single number or (sH, sW). Default: 1

padding Single number or (padH, padW). Default: 0

dilation Single number or (dH, dW). Default: 1

groups split input into groups. Default: 1

Padding



Jonathan Hui

If we pad the input with zeros then the input and output can have the same spatial dimensions.

Zero Padding in NumPy

In NumPy we can add a zero padding of width p to an image as follows:

```
padded = np.zeros(W + 2*p, H + 2*p)  
  
padded[p:W+p, p:H+p] = x
```

Padding


$$L'_\ell = \text{Padd}(L_\ell, p)$$

$$L_{\ell+1}[b, x, y, j] =$$

$$\sigma \left(\left(\sum_{\Delta x, \Delta y, i} W[\Delta x, \Delta y, i, j] L'_\ell[b, x + \Delta x, y + \Delta y, i] \right) - B[j] \right)$$

If the input is padded but the output is not padded then Δx and Δy are non-negative.

Reducing Spatial Dimension



Strides

We can move the filter by a “stride” s for each spatial step.

$$L_{\ell+1}[b, \textcolor{red}{x}, \textcolor{red}{y}, j] = \sigma \left(\left(\sum_{\Delta x, \Delta y, i} W[\Delta x, \Delta y, i, j] L_{\ell}[b, \textcolor{red}{s} * \textcolor{red}{x} + \Delta x, \textcolor{red}{s} * \textcolor{red}{y} + \Delta y, i] \right) - B[j] \right)$$


For strides greater than 1 the spatial dimension is reduced.

Max Pooling

$$L_{\ell+1}[b, \textcolor{red}{x}, \textcolor{red}{y}, i] = \max_{\Delta x, \Delta y} L_{\ell}[b, \textcolor{red}{s} * \textcolor{red}{x} + \Delta x, \textcolor{red}{s} * \textcolor{red}{y} + \Delta y, i]$$

This is typically done with a stride greater than one so that the image dimension is reduced.

Fully Connected (FC) Layers



Fully Connected (FC) Layers

We reshape $L_\ell[b, x, y, i]$ to $L_\ell[b, i']$ and then

$$L_{\ell+1}[b, j] = \sigma \left(\left(\sum_{i'} W[j, i'] L_\ell[b, i'] \right) - B[j] \right)$$

Image to Column (Im2C)

Reduce convolution to matrix multiplication
more space but faster.

$$\begin{aligned} & \tilde{L}_{\ell+1}[b, x, y, j] \\ &= \left(\sum_{\Delta x, \Delta y, i} W[\Delta x, \Delta y, i, j] * L_{\ell}[b, x + \Delta x, y + \Delta y, i] \right) + B[j] \end{aligned}$$

We make a bigger tensor \tilde{L} with two additional indices.

$$\tilde{L}_{\ell}[b, x, y, \Delta x, \Delta y, i] = L_{\ell}[b, x + \Delta x, y + \Delta y, i]$$

Image to Column (Im2C)

$$\tilde{L}_{\ell+1}[b, x, y, j]$$

$$= \left(\sum_{\Delta x, \Delta y, i} W[\Delta x, \Delta y, i, j] * L_{\ell}[b, x + \Delta x, y + \Delta y, i] \right) + B[j]$$


$$= \left(\sum_{\Delta x, \Delta y, i} \tilde{L}_{\ell}[b, x, y, \Delta x, \Delta y, i] * W[\Delta x, \Delta y, i, j] \right) + B[j]$$

$$= \left(\sum_{(\Delta x, \Delta y, i)} \tilde{L}_{\ell}[(b, x, y), (\Delta x, \Delta y, i)] * W[(\Delta x, \Delta y, i), j] \right) + B[j]$$

Dilation


A CNN for image classification typically reduces an $N \times N$ image to a single feature vector.

Dilation is a trick for treating the whole CNN as a “filter” that can be passed over an $M \times M$ image with $M > N$.



An output tensor with full spatial dimension can be useful in, for example, image segmentation.

Dilation



This is called a “fully convolutional” CNN.

Dilation

To implement a fully convolutional CNN we can “dilate” the filters by a dilation parameter d .

$$\tilde{L}_{\ell+1}[b, x, y, j] = W[\Delta x, \Delta y, i, j] L_{\ell}[b, x + d * \Delta x, y + d * \Delta y, i] + B[j]$$

Hypercolumns

An alternative to dilation and fully convolutional networks is hypercolumns.

$$L_{\ell+1}[b, x, y] = L_1[b, x, y]; \cdots ; L_n[b, x/W_\ell, y/W_\ell]; \cdots ; L_N[b]$$

where $x;y$ denotes the concatenation of vectors x and y .

Grouping

$$L_{\ell+1}[b, x, y] = L_{\ell+1}^1[b, x, y]; \cdots ; L_{\ell+1}^G[b, x, y]$$

$$\tilde{L}_{\ell+1}^g[b, x, y, j] = \left(\sum_{\Delta x, \Delta y, i} W^g[\Delta x, \Delta y, i, j] * L_{\ell}[b, x + \Delta x, y + \Delta y, i] \right)$$

For a fixed number of features j in the total output $L_{\ell+1}[b, x, y, j]$ using G groups reduces the number of weight parameters by a factor of G .

2D CNN in PyTorch

conv2d(input, weight, bias, stride, padding, dilation, groups)

input tensor (minibatch,in-channels,iH,iW)

weight filters (out-channels, in-channels/groups, in-channels,kH,kW)

bias tensor (out-channels) . Default: None

stride Single number or (sH, sW). Default: 1

padding Single number or (padH, padW). Default: 0

dilation Single number or (dH, dW). Default: 1

groups split input into groups. Default: 1

Modern Trends

Modern Convolutions use 3X3 filters. This is faster and has fewer parameters. Expressive power is preserved by increasing depth with many stride 1 layers.

Max pooling and dilation seem to have disappeared.

Resnet and resnet-like architectures are now dominant (next lecture).

Alexnet

Given Input[227, 227, 3]

$$L_1[55 \times 55 \times 96] = \text{ReLU}(\text{CONV}(\text{Input}, \Phi_1, \text{width } 11, \text{pad } 0, \text{stride } 4))$$

$$L_2[27 \times 27 \times 96] = \text{MaxPool}(L_1, \text{width } 3, \text{stride } 2))$$

$$L_3[27 \times 27 \times 256] = \text{ReLU}(\text{CONV}(L_2, \Phi_3, \text{width } 5, \text{pad } 2, \text{stride } 1))$$

$$L_4[13 \times 13 \times 256] = \text{MaxPool}(L_3, \text{width } 3, \text{stride } 2))$$

$$L_5[13 \times 13 \times 384] = \text{ReLU}(\text{CONV}(L_4, \Phi_5, \text{width } 3, \text{pad } 1, \text{stride } 1))$$

$$L_6[13 \times 13 \times 384] = \text{ReLU}(\text{CONV}(L_5, \Phi_6, \text{width } 3, \text{pad } 1, \text{stride } 1))$$

$$L_7[13 \times 13 \times 256] = \text{ReLU}(\text{CONV}(L_6, \Phi_7, \text{width } 3, \text{pad } 1, \text{stride } 1))$$


$$L_8[6 \times 6 \times 256] = \text{MaxPool}(L_7, \text{width } 3, \text{stride } 2))$$

$$L_9[4096] = \text{ReLU}(\text{FC}(L_8, \Phi_9))$$

$$L_{10}[4096] = \text{ReLU}(\text{FC}(L_9, \Phi_{10}))$$

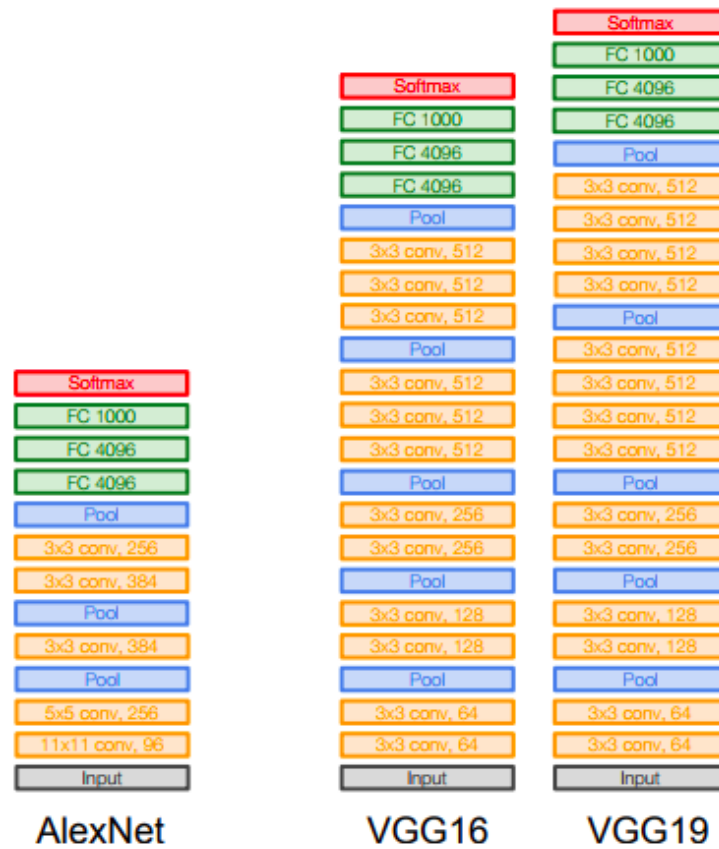
$$s[1000] = \text{ReLU}(\text{FC}(L_{10}, \Phi_s)) \quad \text{class scores}$$

VGG, Zisserman, 2014




Davi Frossard

VGG



Stanford CS231

Inception, Google, 2014




Models for Image Classification in PyTorch

- AlexNet
- VGG
- ResNet
- SqueezeNet
- DenseNet
- Inception v3
- GoogLeNet
- ShuffleNet v2
- MobileNet v2
- ResNeXt
- Wide ResNet
- MNASNet

Imagenet Classification

1000 kinds of objects.



(slide from Kaiming He's recent presentation)

2016 is 3.0%, is 2017 2.25%

SOTA as of January 2020 is 1.3%

END